

ASSIGNMENT-02

DATA VISUALIZATION AND PRE PROCESSING

Assignment Date	22 September 2022
Student Name	SAHANA J M
Student Roll Number	113219071033
Maximum Marks	2 Marks

1. Download the dataset: Dataset

Dataset downloaded in csv form – Churn_Modelling.csv

2. Load the dataset.

```
import pandas as pd
df = pd.read_csv("/content/drive/MyDrive/IBM Assignments/Churn_Modelling.csv")
```

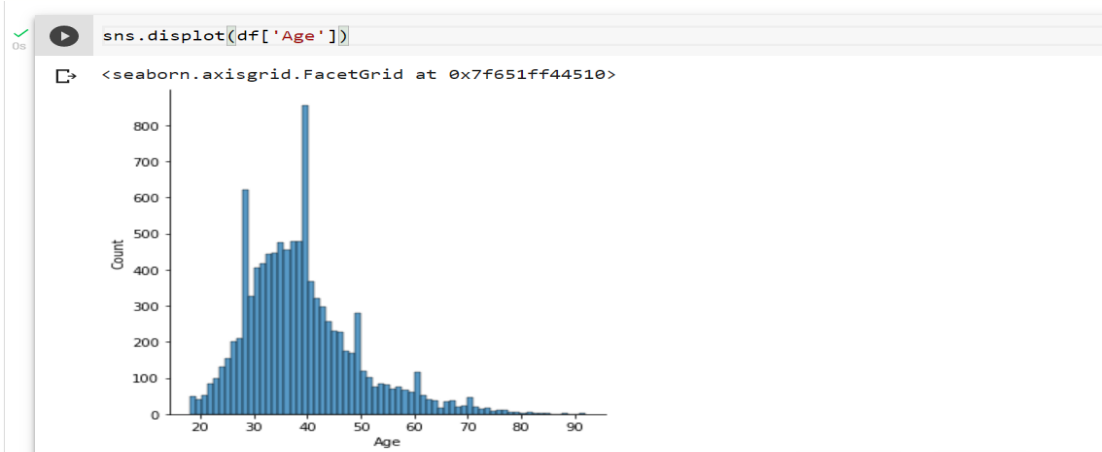
```
✓ 1s import pandas as pd
df = pd.read_csv("/content/drive/MyDrive/IBM Assignments/Churn_Modelling.csv")
```

3. Perform Below Visualizations.


- Univariate Analysis – deals with single column

```
✓ 1s [2] import matplotlib.pyplot as plt
      %matplotlib inline
      import seaborn as sns
```

```
sns.displot(df['Age'])
```

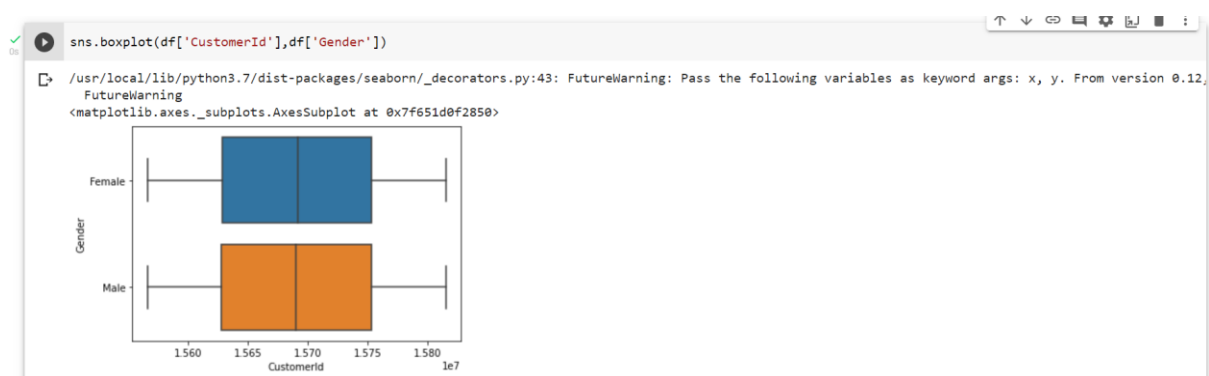


```
sns.displot(df['Age'], kind="kde")
```

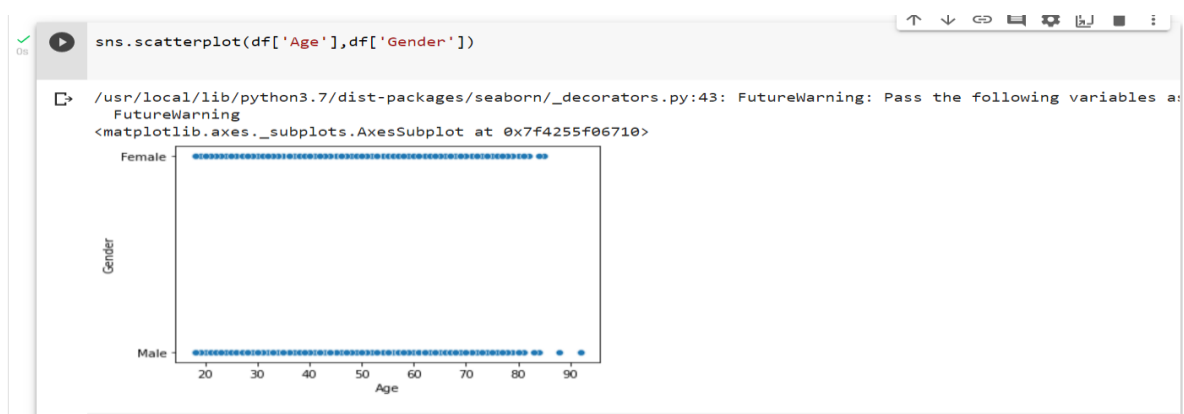


- Bi - Variate Analysis – deals with 2 columns of data


```
sns.boxplot(df['CustomerId'], df['Gender'])
```



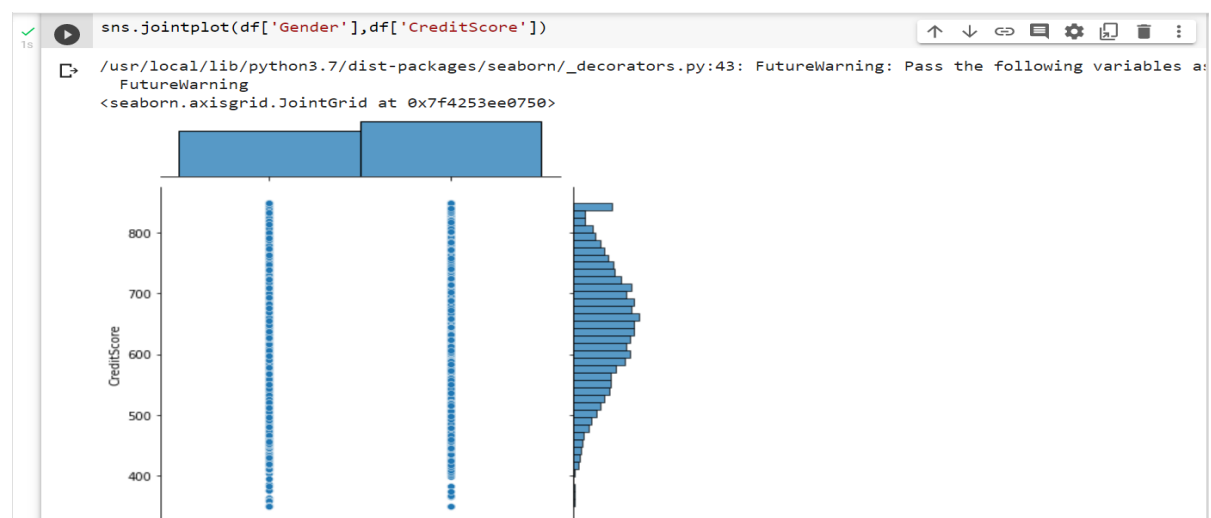
```
sns.scatterplot(df['Age'], df['Gender'])
```



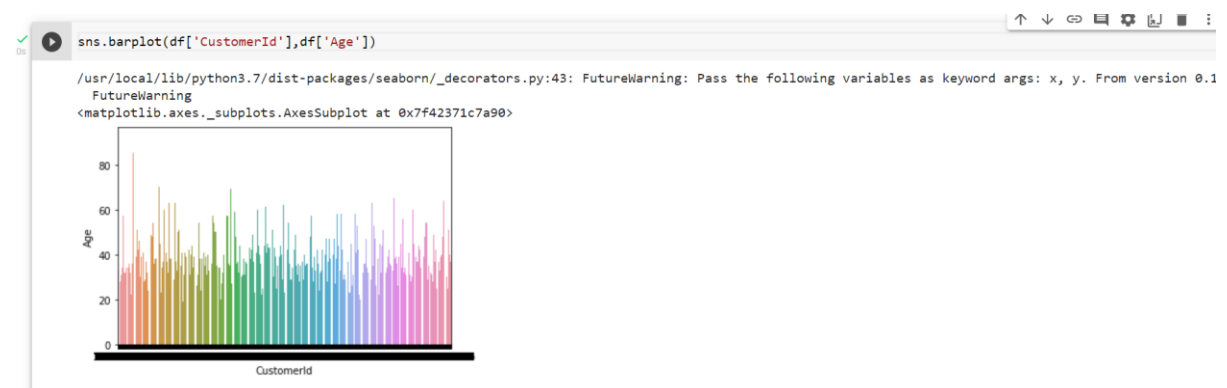
```
sns.scatterplot(df['Age'],df['EstimatedSalary'])
```



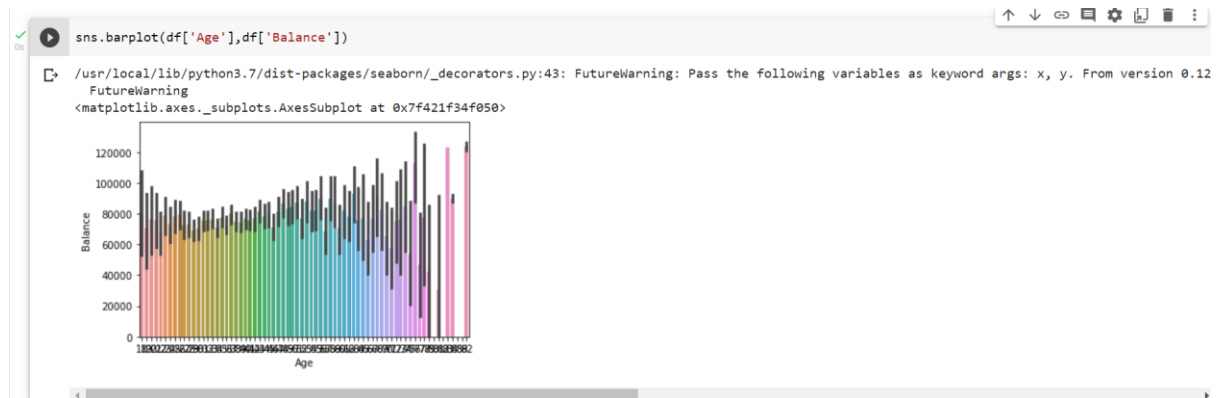
```
sns.jointplot(df['Gender'],df['CreditScore'])
```



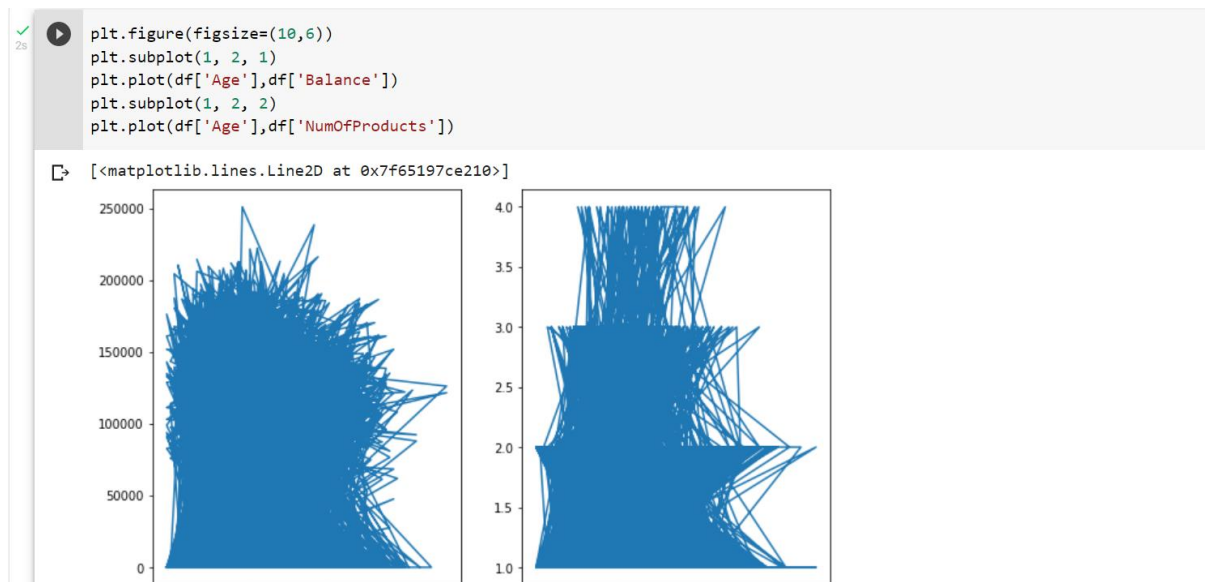
```
sns.barplot(df['CustomerId'],df['Age'])
```



```
sns.barplot(df['Age'],df['Balance'])
```



- Multi - Variate Analysis – deals with multiple columns



4. Perform descriptive statistics on the dataset.

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

Mean: `df.mean()`

```
df.mean()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only')
"""Entry point for launching an IPython kernel.
RowNumber      5.000500e+03
CustomerId      1.569094e+07
CreditScore     6.505288e+02
Age             3.892180e+01
Tenure          5.012800e+00
Balance         7.648589e+04
NumOfProducts  1.530200e+00
HasCrCard       7.055000e-01
IsActiveMember  5.151000e-01
EstimatedSalary 1.000902e+05
Exited          2.037000e-01
dtype: float64
```

Median: `df.median()`

```
df.median()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only')
"""Entry point for launching an IPython kernel.
RowNumber      5.000500e+03
CustomerId      1.569074e+07
CreditScore     6.520000e+02
Age             3.700000e+01
Tenure          5.000000e+00
Balance         9.719854e+04
NumOfProducts  1.000000e+00
HasCrCard       1.000000e+00
IsActiveMember  1.000000e+00
EstimatedSalary 1.001939e+05
Exited          0.000000e+00
dtype: float64
```

Standard Deviation: `df.std()`

```
df.std()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only')
"""Entry point for launching an IPython kernel.
RowNumber      2886.895680
CustomerId     71936.186123
CreditScore     96.653299
Age             10.487806
Tenure          2.892174
Balance        62397.405202
NumOfProducts   0.581654
HasCrCard       0.455840
IsActiveMember   0.499797
EstimatedSalary 57510.492818
Exited          0.402769
dtype: float64
```

5. Handle the Missing values.

`df.isnull().sum()`

```
df.isnull().sum()

RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

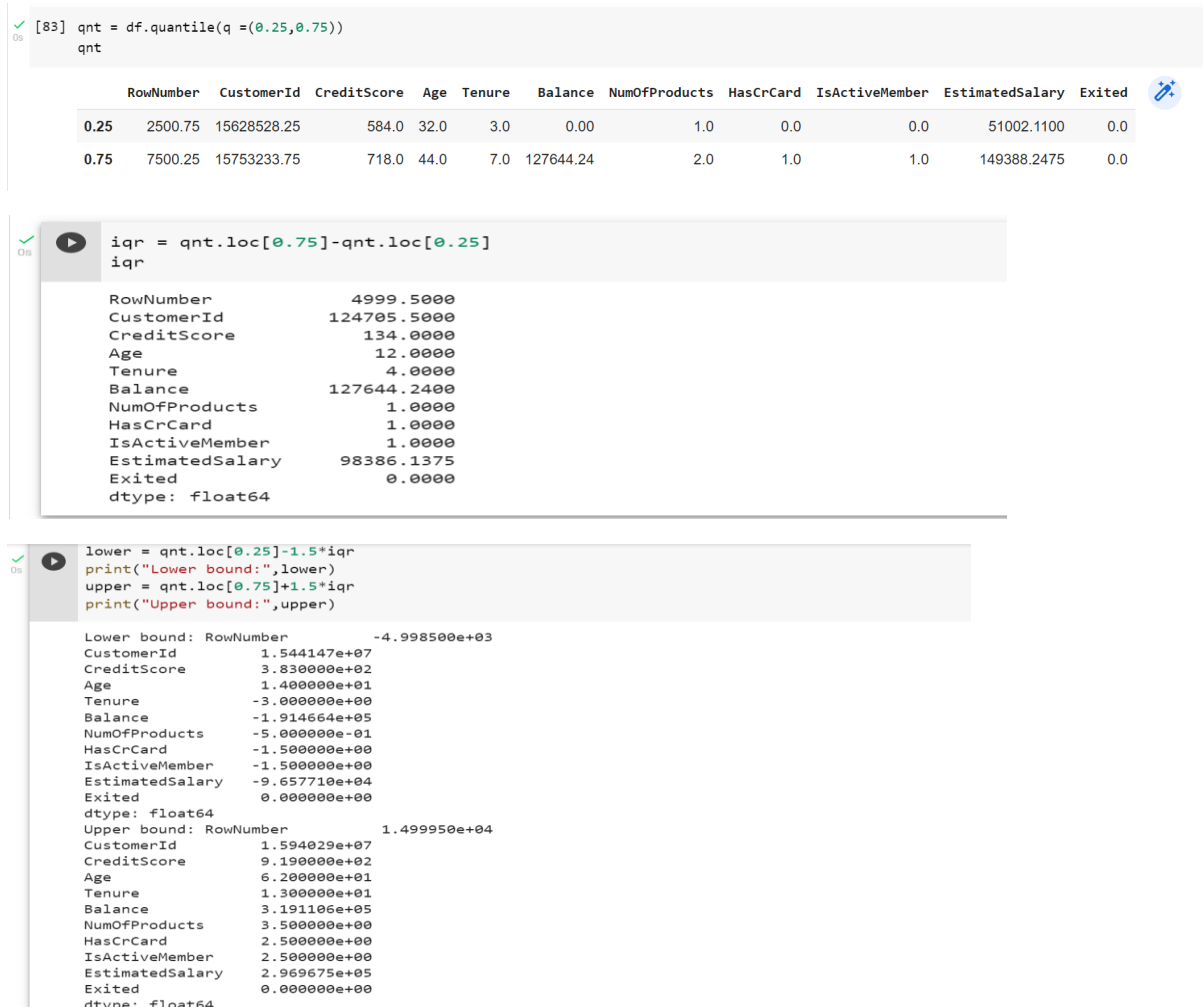
6. Find the outliers and replace the outliers

Outliers: The values which are different from others or less relative to others.

Using Boxplot



Using method



Replacing Outliers:

```
df['Balance'] = np.where(df['Balance']>127644,0.00,df['Balance'])
```

```
''' replacing outliers '''
df['Balance'] = np.where(df['Balance']>127644,0.00,df['Balance'])
```

7. Check for Categorical columns and perform encoding.

Categorical columns: Geography, Gender

Encoding changes the values to numerical forms such as 0,1

```
[98] from sklearn.preprocessing import LabelEncoder
labelencoder_df = LabelEncoder()
df['Geography'] = labelencoder_df.fit_transform(df['Geography'])

df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	0	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	2	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	0	Female	42	8	0.00	3	1	0	113931.57
3	4	15701354	Boni	699	0	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	2	Female	43	2	125510.82	1	1	1	79084.10

```
df['Gender'] = labelencoder_df.fit_transform(df['Gender'])

df.head(7)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	0	0	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	0	0	42	8	0.00	3	1	0	113931.57
3	4	15701354	Boni	699	0	0	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	2	0	43	2	125510.82	1	1	1	79084.10
5	6	15574012	Chu	645	2	1	44	8	113755.78	2	1	0	149756.71
6	7	15592531	Bartlett	822	0	1	50	7	0.00	2	1	1	10062.80

8. Split the data into dependent and independent variables.

```
[105] X = df.iloc[:, :-1].values
print(X)

[[1 15634602 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619304 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42085.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]

Y = df.iloc[:, -1].values
print(Y)

[1 0 1 ... 1 1 0]
```

9. Scale the independent variables

```
[115] from sklearn.preprocessing import scale
      Y = scale(Y)
```

Y

```
array([[ 1.97716468, -0.50577476,  1.97716468, ...,  1.97716468,
         1.97716468, -0.50577476]])
```

10. Split the data into training and testing

Y_train

```
array([-0.50577476, -0.50577476, -0.50577476, ..., -0.50577476,
       -0.50577476,  1.97716468])
```

Y_test

```
array([-0.50577476,  1.97716468, -0.50577476, ..., -0.50577476,
       -0.50577476, -0.50577476])
```

X_train

```
array([[7390, 15676909, 'Mishin', ..., 1, 0, 163830.64],
       [9276, 15749265, 'Carslaw', ..., 1, 1, 57098.0],
       [2996, 15582492, 'Moore', ..., 1, 0, 185630.76],
       ...,
       [3265, 15574372, 'Hoolan', ..., 1, 0, 181429.87],
       [9846, 15664035, 'Parsons', ..., 1, 1, 148750.16],
       [2733, 15592816, 'Udokamma', ..., 1, 0, 118855.26]], dtype=object)
```

X_test

```
array([[9395, 15615753, 'Upchurch', ..., 1, 1, 192852.67],
       [899, 15654700, 'Fallaci', ..., 1, 0, 128702.1],
       [2399, 15633877, 'Morrison', ..., 1, 1, 75732.25],
       ...,
       [9550, 15772604, 'Chiemezie', ..., 1, 0, 141533.19],
       [2741, 15787699, 'Burke', ..., 1, 1, 11276.48],
       [6691, 15579223, 'Niu', ..., 1, 0, 192950.6]], dtype=object)
```